

TRP Week 7 - Data Contract Enforcer (Automatically generated report)

Data Health Score: 38 57 100

Violations (top 3):

- CRITICAL: Week 2 Digital Courtroom: verdict records failed structured LLM output schema validation Downstream AI consumers can ingest malformed scores or evidence and produce incorrect risk decisions
- CRITICAL: Week 3 Document Refinery: extracted facts confidence must stay a float in 0 0-1 0 Downstream systems that interpret confidence as a probability (Week 4 lineage attribution and any scoring consumers) may mis-rank or mis-blame ...
- CRITICAL: Week 3 Document Refinery: extracted facts entity refs must reference IDs present in entities Downstream lineage metadata can become inconsistent, breaking entity-aware attribution and any audit joins

Schema changes detected:

- Week 3 extraction records: Field extracted facts changed type from array to object Breaking change requires coordinated consumer updates
- Week 5 event records: Field payload changed type from object to integer Breaking change requires coordinated consumer updates

AI system risk assessment:

- Embedding drift is within acceptable bounds
- Structured LLM output schema violation rate is rising (violation rate=0 12, baseline=0 0)

Recommended actions:

- 1 Update src week3 extractor py so extracted facts confidence is emitted as a float in the 0 0-1 0 range (contract clause: week3 extracted facts confidence range)
- 2 Update src week3 extractor py so every extracted facts entity refs value is an entity id present in the same record's entities array (contract clause: week3 extracted facts entity refs relationship)
- 3 Stabilize Week 2 structured verdict outputs by updating the LLM prompt or parser so the scores score field is always an integer 1-5 (AI contract: week2 verdict record schema)